


Conviction masking — cora (BundleSheaf, d=5, 3 layers)
avg z-scored conviction (layers 1..L) | feature masking (train pool only) | full test set | mean $\pm$ SE, 10 folds

Accuracy on unmasked test nodes



Mean confidence on unmasked test nodes

