


Conviction hard message gating — cora (GeneralSheaf, d=5, 3 layers)
avg z-scored conviction (layers 1..L) | hard message gating (train pool only) | full test set | mean $\pm$ SE, 10 folds

Accuracy on test set



Mean confidence on test set

