


Conviction hard message gating — cornell (GeneralSheaf, d=2, 2 layers)  
avg z-scored conviction (layers 1..L) | hard message gating (train pool only) | full test set | mean  $\pm$  SE, 10 folds

**Accuracy on test set**



**Mean confidence on test set**

