


Conviction hard message gating — cornell (GeneralSheaf, d=2, 2 layers)
X_L conviction | hard message gating (train pool only) | full test set | mean $\pm$ SE, 10 folds

Accuracy on test set



Mean confidence on test set

