


Conviction soft gating — amazon_ratings (GeneralSheaf, d=5, 5 layers)
avg z-scored conviction (layers 1..L) | entire train pool gated | full test set | mean $\pm$ SE, 10 folds

Accuracy on test set



Mean confidence on test set

