


Conviction feature masking — amazon_ratings (GeneralSheaf, d=5, 5 layers)
avg z-scored conviction (layers 1..L) | feature masking (train pool only) | full test set | mean $\pm$ SE, 10 folds

Accuracy on test set



Mean confidence on test set

