


Conviction retrain — hard message gating — cornell (GeneralSheaf, d=2, 2 layers)  
X\_L conviction | hard message gating, retrain-from-scratch (train pool only) | full test set | mean  $\pm$  SE, 10 folds

**Accuracy on test set**



**Mean confidence on test set**

