


Conviction retrain — feature masking — roman_empire (GeneralSheaf, d=5, 5 layers)
X_L conviction | feature masking, retrain-from-scratch (train pool only) | full test set | mean $\pm$ SE, 10 folds

Accuracy on test set



Mean confidence on test set

