


Conviction masking — cora (GeneralSheaf, d=5, 3 layers)
X_L conviction | feature masking (train+val pool) | full test set | mean $\pm$ SE, 10 folds

Accuracy on unmasked test nodes



Mean confidence on unmasked test nodes

